

Supervised Stochastic Triplet Embedding

November 10, 2015

1 Projected Gradient

After we calculate the gradient using a loss function (in this case, the loss for stochastic triplet embedding, STE), we get an $n \times n$ matrix \mathbf{K} . This matrix needs to be positive semidefinite (having no negative eigenvalues) in order to be a valid kernel. STE uses projected gradient descent - after each gradient step at iteration t , where we update the variable \mathbf{K}^t , we then project \mathbf{K} to the closest matrix in the set of positive semidefinite matrices, which is a closed convex set (a cone), giving the following optimization problem

$$\begin{aligned} & \text{minimize} && \rho \|\mathbf{K} - \mathbf{X}\|_F^2 \\ & \text{s.t} && \mathbf{X} \succeq 0, \end{aligned} \tag{1}$$

which we can solve by calculating the eigenvalue decomposition of \mathbf{K} and setting all negative eigenvalues to zero.

In metric learning, we want to impose a stronger restriction on \mathbf{K} - instead of being simply psd, we additionally require that the dot products are a function of some features, $\mathbf{K} = \mathbf{F}^T \mathbf{M} \mathbf{F}$. There are many algorithms for metric learning, but it's simple to adapt projected gradient to this problem. Given our feature matrix \mathbf{F} , we solve a very similar optimization problem,

$$\begin{aligned} & \text{minimize} && \rho \|\mathbf{X} - \mathbf{K}\|_F^2 \\ & \text{s.t} && \mathbf{X} = \mathbf{F}^T \mathbf{M} \mathbf{F} \\ & && \mathbf{M} \succeq 0 \end{aligned} \tag{2}$$

which we can solve in almost as simple a way, after some linear algebra. Below we derive the updates for the objective function, and then we can again require that \mathbf{M} be psd, again by projection.

$$\|\mathbf{K} - \mathbf{F}^T \mathbf{M} \mathbf{F}\|_F^2 = \tag{3}$$

$$\text{Tr}((\mathbf{K} - \mathbf{F}^T \mathbf{M} \mathbf{F})^T (\mathbf{K} - \mathbf{F}^T \mathbf{M} \mathbf{F})) \tag{4}$$

$$\text{Tr}(\mathbf{K}^T \mathbf{K} - \mathbf{F}^T \mathbf{M} \mathbf{F} \mathbf{K} - \mathbf{K}^T \mathbf{F}^T \mathbf{M} \mathbf{F} + \mathbf{F}^T \mathbf{M} \mathbf{F} \mathbf{F}^T \mathbf{M} \mathbf{F}) \tag{5}$$

$$\text{Tr}(\mathbf{K}^T \mathbf{K}) - 2 \text{Tr}(\mathbf{F}^T \mathbf{M} \mathbf{F} \mathbf{K}) + \text{Tr}(\mathbf{F}^T \mathbf{M} \mathbf{F} \mathbf{F}^T \mathbf{M} \mathbf{F}) \tag{6}$$

where we use a few properties of trace and the fact that all the relevant matrices are symmetric.

We can solve this gradient analytically, first by taking derivatives and then setting them to zero

$$-2\mathbf{F}^T\mathbf{K}\mathbf{F} + 2\mathbf{F}^T\mathbf{F}\mathbf{M}\mathbf{F}^T\mathbf{F} = 0 \quad (7)$$

$$\mathbf{F}^T\mathbf{K}\mathbf{F} = \mathbf{F}^T\mathbf{F}\mathbf{M}\mathbf{F}^T\mathbf{F} \quad (8)$$

$$(9)$$

Then we multiply both sides by the inverse of $(\mathbf{F}^T\mathbf{F})$ ¹.

$$\mathbf{M} = (\mathbf{F}^T\mathbf{F})^{-1}\mathbf{F}^T\mathbf{K}\mathbf{F}(\mathbf{F}^T\mathbf{F})^{-1} \quad (10)$$

Note that this is the sample covariance, and the inverse is sometimes known as the precision matrix. The update is just a linear function of \mathbf{K} , and in general we can't expect \mathbf{K} to be positive semidefinite, so we still need to make sure that \mathbf{M} is positive semidefinite. Note that the main work in this step is inversion of $\mathbf{F}^T\mathbf{F}$, and this never changes as we iterate, so this can be cached.

2 Results

In this section I compare supervised STE with unsupervised STE. Unsurprisingly, it does much better when it has features, because there are fewer parameters to learn: instead of coordinates for each point in a possibly high dimensional space, it has to learn how to rotate and scale the features - the size of the matrix \mathbf{M} only depends on the number of features, not the number of data points.

¹TODO : check this again with the pseudoinverse to make sure it's still true

numTriplets

supervised

unsupervised

10

100

1000

10000

